

K-Means Clustering of Land Use Patterns in London

Project Goal:

This project aims to classify London's wards into spatially distinct groups based on detailed land use compositions using the K-Means clustering algorithm. By analyzing the percentage share of key land types—such as domestic buildings, greenspace, roads, and infrastructure—six distinct urban typologies are revealed, ranging from dense commercial cores to suburban residential zones and green corridors. These clusters are intended to provide valuable spatial insights for urban planning, land use policy, and neighborhood characterization across Greater London.

Data Source: London Datastore – Greater London Authority.

- [Land Use by Borough and Ward](#)
- [2011 London Output Area Classification](#)

1. Load and Explore the Dataset

```
In [1]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import seaborn as sns
import os
import warnings
import geopandas as gpd

# Optional: Suppress specific warnings
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=UserWarning)

# Optional: Limit MKL threads to avoid memory warnings (Windows only)
os.environ["OMP_NUM_THREADS"] = "3"
```

```
In [2]: # Load the dataset
df = pd.read_csv("land-use-glud-ward.csv")
```

```
In [3]: # Display initial information
df.head()
```

Out[3]:

	Area Code	Local Authority Name	Area name	Area of Domestic Buildings	Area of Domestic Gardens	Area of Non Domestic Buildings	Area of Road	Area of Rail	Area of Path	Area of Greenspace	...	% Area of Non Domestic Buildings	% Area of Road	% Area of Rail	% Area of Path	% Area of Greenspace	% Area of Water
0	00AA	City of London	City of London	253.49	4.02	1254.69	797.01	7.53	13.58	161.75	...	37.5	23.8	0.2	0.4	4.8	13.3
1	00ABFX	Barking and Dagenham	Abbey	149.04	198.92	112.17	243.22	84.44	28.06	255.31	...	8.7	18.9	6.6	2.2	19.8	2.0
2	00ABFY	Barking and Dagenham	Alibon	153.87	575.82	37.20	184.36	27.01	5.29	311.80	...	2.7	13.5	2.0	0.4	22.9	0.0
3	00ABFZ	Barking and Dagenham	Becontree	193.96	626.05	43.68	260.83	0	7.31	74.92	...	3.4	20.5	0.0	0.6	5.9	0.0
4	00ABGA	Barking and Dagenham	Chadwell Heath	205.58	602.38	64.79	315.29	0	26.09	2505.92	...	1.7	8.1	0.0	0.7	64.3	0.3

5 rows × 27 columns

```
In [4]: # Summary of the DataFrame
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 626 entries, 0 to 625
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Area Code        626 non-null    object  
 1   Local Authority Name 625 non-null    object  
 2   Area name        626 non-null    object  
 3   Area of Domestic Buildings 626 non-null    float64 
 4   Area of Domestic Gardens   626 non-null    float64 
 5   Area of Non Domestic Buildings 626 non-null    float64 
 6   Area of Road       626 non-null    float64 
 7   Area of Rail       626 non-null    object  
 8   Area of Path       626 non-null    float64 
 9   Area of Greenspace 626 non-null    float64 
 10  Area of Water      626 non-null    float64 
 11  Area of Other Land Uses 626 non-null    float64 
 12  Area of Unclassified Land 626 non-null    float64 
 13  Total Area of All Land Types 626 non-null    float64 
 14  Area of Admin Geography 626 non-null    float64 
 15  % Area of Domestic Buildings 626 non-null    float64 
 16  % Area of Domestic Gardens   626 non-null    float64 
 17  % Area of Non Domestic Buildings 626 non-null    float64 
 18  % Area of Road       626 non-null    float64 
 19  % Area of Rail       625 non-null    float64 
 20  % Area of Path       626 non-null    float64 
 21  % Area of Greenspace 626 non-null    float64 
 22  % Area of Water      625 non-null    float64 
 23  % Area of Other Land Uses 626 non-null    float64 
 24  % Area of Unclassified Land 626 non-null    float64 
 25  % Total Area of All Land Types 626 non-null    float64 
 26  Quality of Fit Indicator 624 non-null    float64 
dtypes: float64(23), object(4)
memory usage: 132.2+ KB
```

- The dataset contains **626 rows** and **27 columns**.
- It includes both absolute values (e.g., "Area of Greenspace") and percentage values (e.g., "% Area of Greenspace") for different land use types.
- Some fields like `Area of Rail` are incorrectly read as `object` types instead of numeric.
- A few missing values exist, mostly in percentage or quality-related fields.

2. Data Cleaning

2.1. Fix Incorrect Data Types

```
In [5]: # Convert 'Area of Rail' from object to float
df["Area of Rail"] = pd.to_numeric(df["Area of Rail"], errors = "coerce")
df.dtypes
```

```
Out[5]: Area Code          object
Local Authority Name  object
Area name            object
Area of Domestic Buildings  float64
Area of Domestic Gardens   float64
Area of Non Domestic Buildings  float64
Area of Road          float64
Area of Rail           float64
Area of Path           float64
Area of Greenspace     float64
Area of Water          float64
Area of Other Land Uses float64
Area of Unclassified Land  float64
Total Area of All Land Types  float64
Area of Admin Geography float64
% Area of Domestic Buildings  float64
% Area of Domestic Gardens   float64
% Area of Non Domestic Buildings  float64
% Area of Road          float64
% Area of Rail           float64
% Area of Path           float64
% Area of Greenspace     float64
% Area of Water          float64
% Area of Other Land Uses float64
% Area of Unclassified Land  float64
% Total Area of All Land Types  float64
Quality of Fit Indicator  float64
dtype: object
```

2.2 Check for Missing Data

```
In [6]: df.isnull().sum()
```

```
Out[6]: Area Code          0  
Local Authority Name    1  
Area name              0  
Area of Domestic Buildings 0  
Area of Domestic Gardens 0  
Area of Non Domestic Buildings 0  
Area of Road             0  
Area of Rail              1  
Area of Path              0  
Area of Greenspace        0  
Area of Water              0  
Area of Other Land Uses   0  
Area of Unclassified Land 0  
Total Area of All Land Types 0  
Area of Admin Geography   0  
% Area of Domestic Buildings 0  
% Area of Domestic Gardens 0  
% Area of Non Domestic Buildings 0  
% Area of Road             0  
% Area of Rail              1  
% Area of Path              0  
% Area of Greenspace        0  
% Area of Water              1  
% Area of Other Land Uses   0  
% Area of Unclassified Land 0  
% Total Area of All Land Types 0  
Quality of Fit Indicator   2  
dtype: int64
```

- Only 5 columns contain missing values.
- The columns with missing data include: `Area of Rail`, `% Area of Rail`, `% Area of Water`, `Local Authority Name`, and `Quality of Fit Indicator`.
- Each of these columns has **only 1 or 2 missing values**
- This low number of missing values allows to safely drop these rows without losing meaningful information.

2.3 Drop Missing Values

```
In [7]: df_cleaned = df.dropna()  
df_cleaned.isnull().sum()
```

```
Out[7]: Area Code          0  
Local Authority Name    0  
Area name              0  
Area of Domestic Buildings 0  
Area of Domestic Gardens 0  
Area of Non Domestic Buildings 0  
Area of Road             0  
Area of Rail              0  
Area of Path              0  
Area of Greenspace        0  
Area of Water              0  
Area of Other Land Uses   0  
Area of Unclassified Land 0  
Total Area of All Land Types 0  
Area of Admin Geography   0  
% Area of Domestic Buildings 0  
% Area of Domestic Gardens 0  
% Area of Non Domestic Buildings 0  
% Area of Road             0  
% Area of Rail              0  
% Area of Path              0  
% Area of Greenspace        0  
% Area of Water              0  
% Area of Other Land Uses   0  
% Area of Unclassified Land 0  
% Total Area of All Land Types 0  
Quality of Fit Indicator   0  
dtype: int64
```

```
In [8]: df_cleaned.shape
```

```
Out[8]: (622, 27)
```

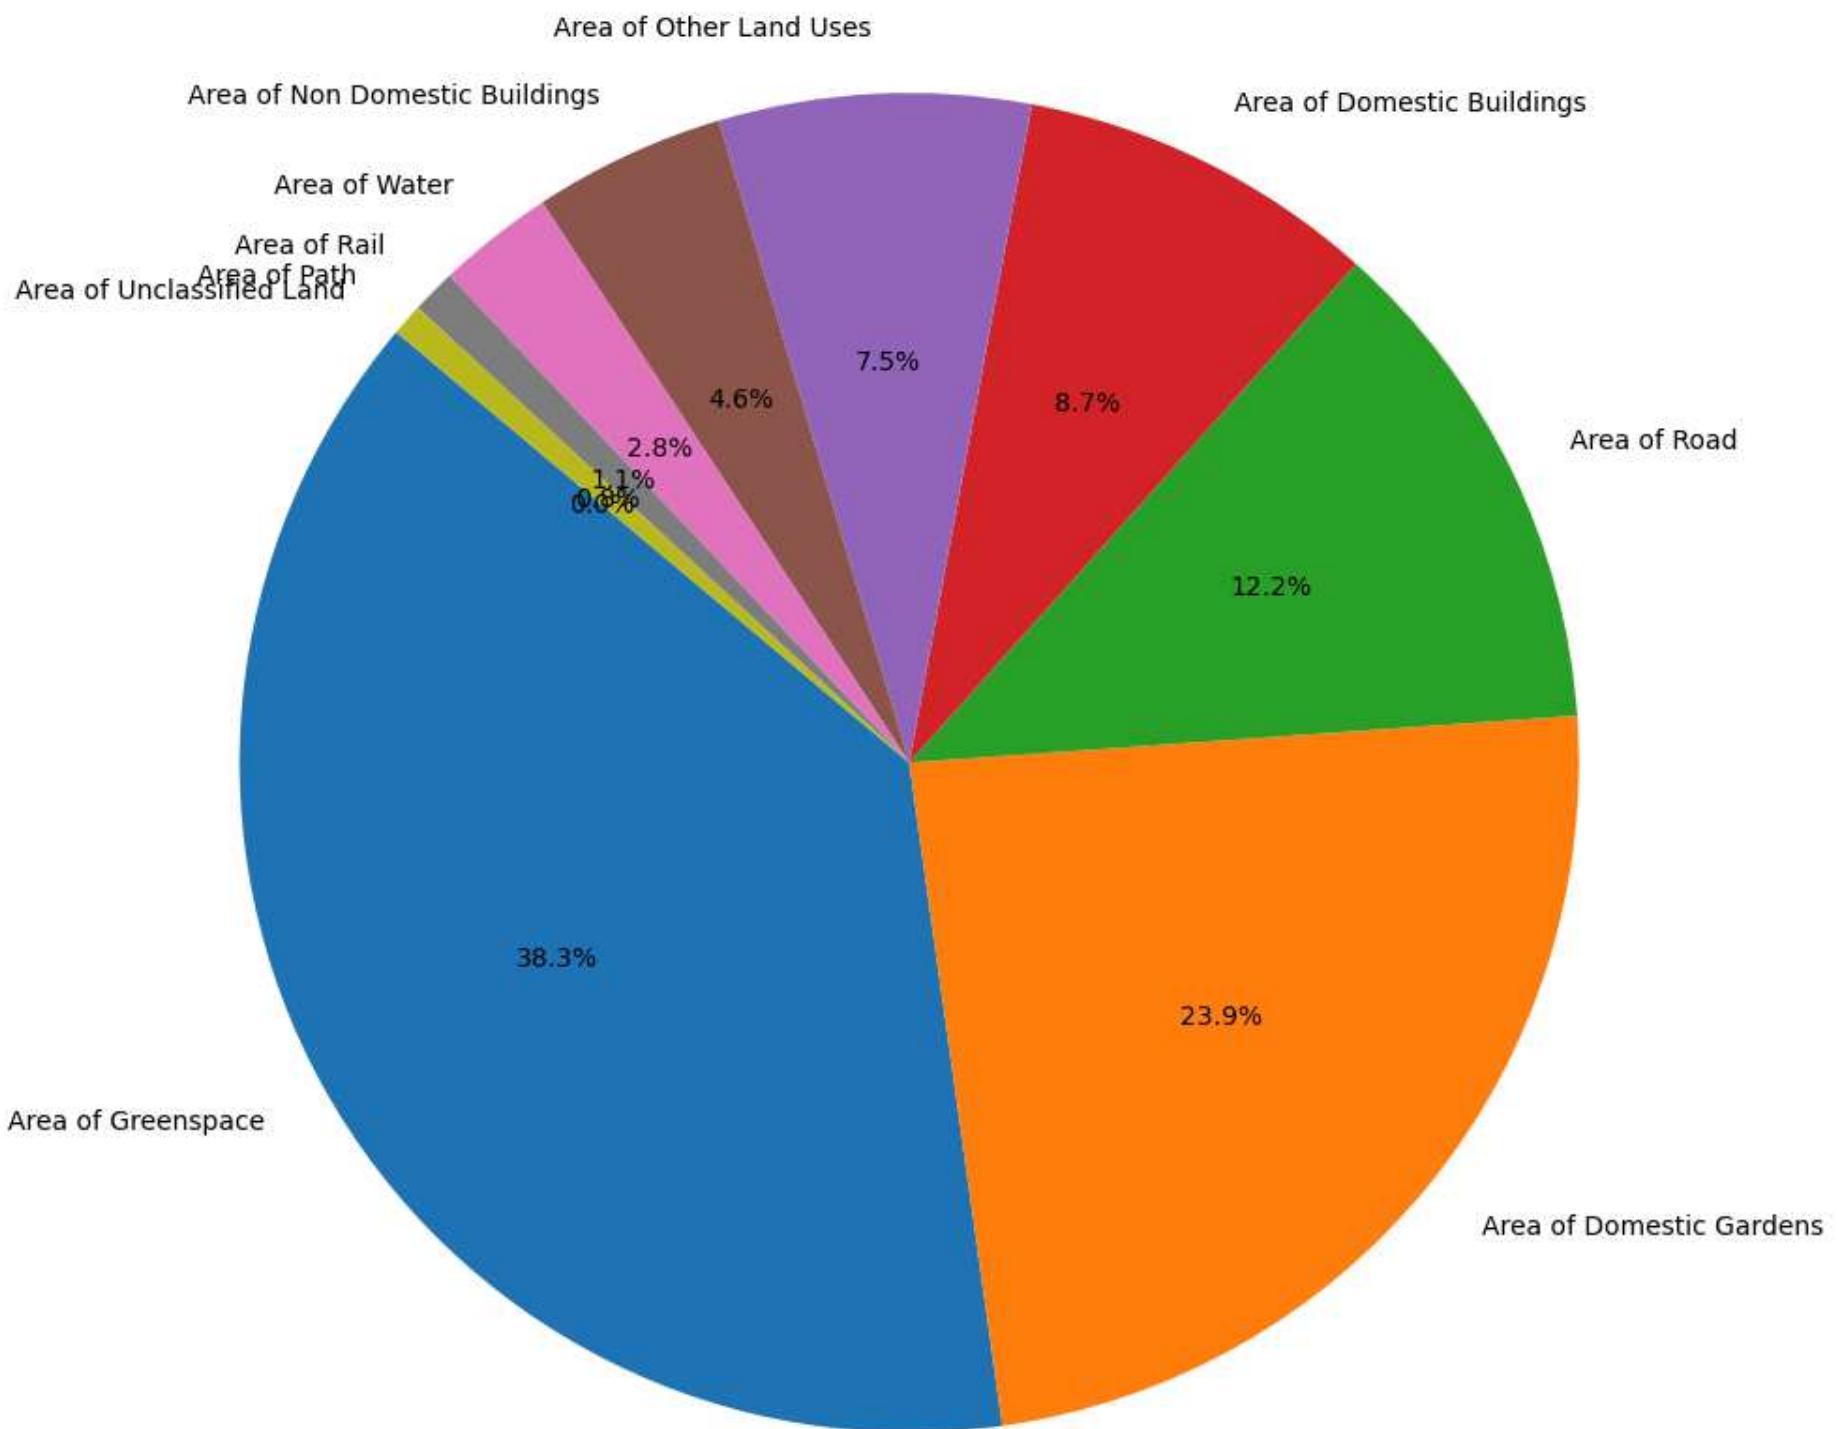
3. Visualize Average Land Use Distribution in London

```
In [9]: land_use_columns = [  
    "Area of Domestic Buildings", "Area of Domestic Gardens",  
    "Area of Non Domestic Buildings", "Area of Road", "Area of Rail",  
    "Area of Path", "Area of Greenspace", "Area of Water",  
    "Area of Other Land Uses", "Area of Unclassified Land"  
]  
mean_land_uses = df_cleaned[land_use_columns].mean().sort_values(ascending=False)
```

```
In [10]: plt.figure(figsize=(10, 10))  
plt.pie(mean_land_uses, labels=mean_land_uses.index, autopct='%1.1f%%', startangle=140)  
plt.title("Average Land Use Distribution in London")  
plt.axis('equal')
```

```
plt.tight_layout()  
plt.show()
```

Average Land Use Distribution in London



- **Greenspace** makes up the largest portion at **38.3%**, indicating a significant amount of natural and recreational space.
- **Domestic Gardens** follow with **23.9%**, reflecting the suburban character of many residential areas.
- **Roads (12.2%)** and **Domestic Buildings (8.7%)** make up much of the built environment.
- Other uses such as **Non-Domestic Buildings, Water, Rail, Paths, and Unclassified Land** account for smaller shares.
- This distribution highlights London's mix of green space and residential land, with moderate urban infrastructure.

4. Zoning with K-Means Clustering

4.1 Why K-Means?

- No labeled zoning categories exist → **unsupervised learning is appropriate**.
- We aim to group similar land use patterns.
- K-Means is a straightforward and effective clustering method for numeric data.

4.2 Prepare Features

```
In [11]: # Select percentage-based Land use columns  
percent_columns = [col for col in df_cleaned.columns if col.startswith('% Area')]  
df_cluster = df_cleaned[percent_columns].dropna()
```

Only percentage-based land use columns were selected in the K-Means clustering analysis to ensure size-invariant comparisons across wards. This selection enabled the detection of meaningful patterns in how land is proportionally allocated, rather than clustering based on the total size of each ward. This approach was adopted to promote interpretability and fairness in the clustering outcomes.

4.3. Feature Scaling

```
In [12]: # Standardize the data for clustering
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df_cleaned[percent_columns])
```

```
In [13]: X_scaled_df = pd.DataFrame(X_scaled, columns=percent_columns)
X_scaled_df.head()
```

Out[13]:

	% Area of Domestic Buildings	% Area of Domestic Gardens	% Area of Non Domestic Buildings	% Area of Road	% Area of Rail	% Area of Path	% Area of Greenspace	% Area of Water	% Area of Other Land Uses	% Area of Unclassified Land
0	-0.060192	-0.900690	0.396700	0.500340	2.132175	1.937470	-0.393314	0.016826	1.189230	-0.040129
1	-0.118329	1.324847	-0.684973	-0.456852	0.252529	-0.856439	-0.206473	-0.313854	-0.740374	-0.040129
2	0.637442	1.889536	-0.558777	0.783952	-0.564708	-0.546005	-1.231084	-0.313854	-0.623428	-0.040129
3	-1.281053	-0.900690	-0.865251	-1.414043	-0.564708	-0.390788	2.288757	-0.264252	-0.857320	-0.040129
4	-1.106645	-0.610041	-0.306387	-1.360866	-0.196952	0.385298	1.173739	0.463244	0.097737	-0.040129

Standardization is important because K-Means is sensitive to scale. Therefore, z-score normalization (mean = 0, standard deviation = 1) was applied to ensure fair comparisons across features.

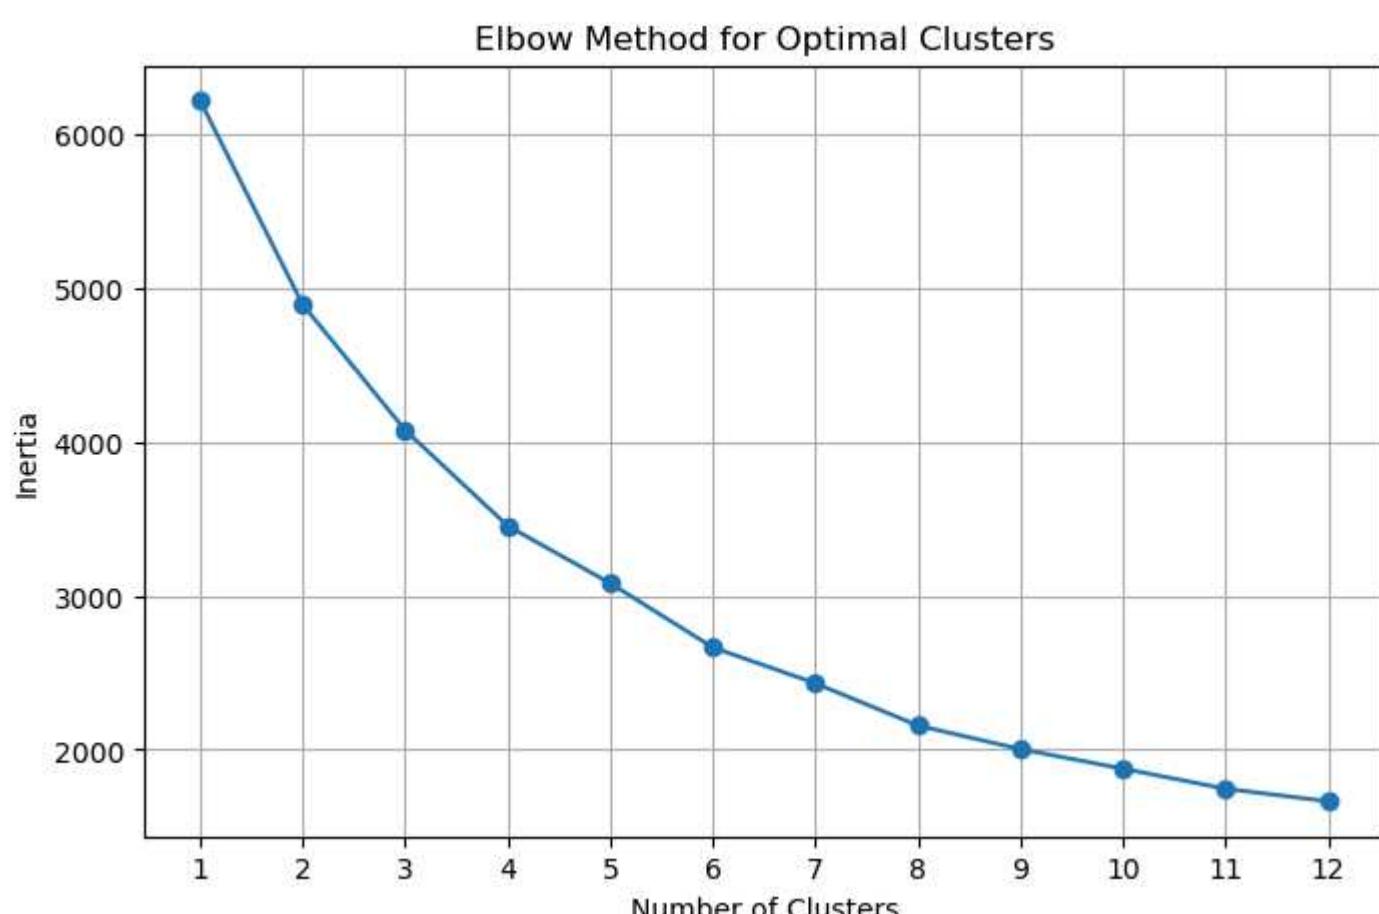
4.4. Determine Optimal Number of Clusters (Elbow Method)

```
In [14]: # Ensure clean slate
inertias = []
k_values = list(range(1, 13))
```

```
In [15]: # Compute inertia for each k
for k in k_values:
    model = KMeans(n_clusters=k, random_state=42, n_init=10)
    model.fit(X_scaled)
    inertias.append(model.inertia_)
# Check that both lists have the same length
print(f"k_values: {len(k_values)} | inertias: {len(inertias)}")
```

```
k_values: 12 | inertias: 12
```

```
In [16]: # Plot the elbow curve
plt.figure(figsize=(8, 5))
plt.plot(k_values, inertias, marker='o')
plt.title('Elbow Method for Optimal Clusters')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.grid(True)
plt.xticks(k_values)
plt.show()
```



Based on the elbow method, the optimal number of clusters appears to be $k = 4$, as this is the point where the inertia begins to level off. However, to capture more nuanced patterns within the data and allow for a finer segmentation of clusters, proceeding with $k = 6$ was chosen as the preferred approach. This decision aims to provide a more detailed understanding of the underlying structure, even if it introduces slightly more complexity.

4.5. Apply K-Means Clustering (k=6)

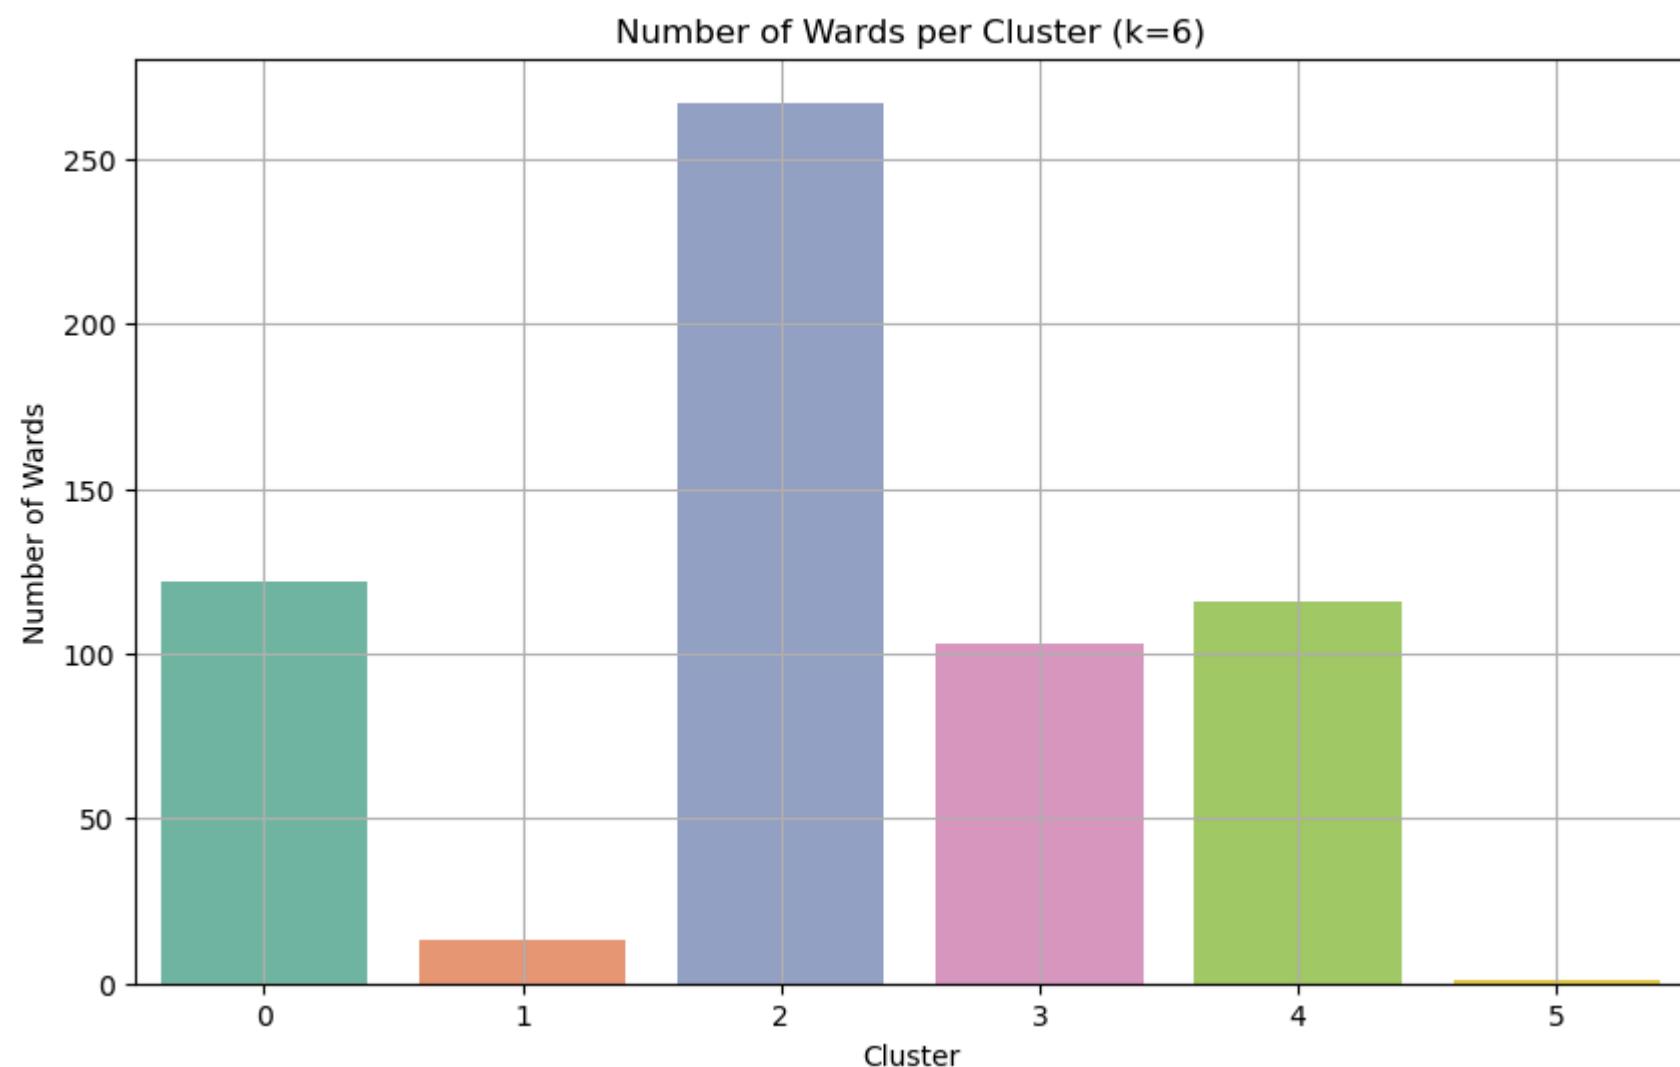
```
In [17]: # Apply KMeans with k=6
kmeans_6 = KMeans(n_clusters=6, random_state=42, n_init=10)
cluster_labels_6 = kmeans_6.fit_predict(X_scaled)
```

```
In [18]: # Assign cluster labels to a new DataFrame
df_with_labels_6 = df_cleaned.loc[df_cluster.index].copy()
df_with_labels_6['Cluster'] = cluster_labels_6
```

```
In [19]: # Cluster sizes
cluster_counts_6 = df_with_labels_6['Cluster'].value_counts().sort_index()
```

```
In [20]: # Plot the distribution
plt.figure(figsize=(10, 6))
sns.barplot(x=cluster_counts_6.index, y=cluster_counts_6.values, palette='Set2')
plt.title('Number of Wards per Cluster (k=6)')
plt.xlabel('Cluster')
plt.ylabel('Number of Wards')
plt.grid(True)
plt.show()

print(cluster_counts_6)
```



```
Cluster
0    122
1     13
2    267
3    103
4    116
5      1
Name: count, dtype: int64
```

With 6 clusters, a moderate level of segmentation is achieved—more detailed than 4 clusters but less fragmented than 10. This number is considered a suitable balance for understanding key urban types in London. Notably, **Cluster 2 was identified as the most populated group**, representing a large number of wards with balanced, mixed-use land compositions.

```
In [21]: df_with_labels_6.head()
```

Out[21]:

	Area Code	Local Authority Name	Area name	Area of Domestic Buildings	Area of Domestic Gardens	Area of Non Domestic Buildings	Area of Road	Area of Rail	Area of Path	Area of Greenspace	...	% Area of Road	% Area of Rail	% Area of Path	% Area of Greenspace	% Area of Water	% Area of Other Land Uses
1	00ABFX	Barking and Dagenham	Abbey	149.04	198.92	112.17	243.22	84.44	28.06	255.31	...	18.9	6.6	2.2	19.8	2.0	14.7
2	00ABFY	Barking and Dagenham	Alibon	153.87	575.82	37.20	184.36	27.01	5.29	311.80	...	13.5	2.0	0.4	22.9	0.0	4.8
3	00ABFZ	Barking and Dagenham	Becontree	193.96	626.05	43.68	260.83	0.00	7.31	74.92	...	20.5	0.0	0.6	5.9	0.0	5.4
4	00ABGA	Barking and Dagenham	Chadwell Heath	205.58	602.38	64.79	315.29	0.00	26.09	2505.92	...	8.1	0.0	0.7	64.3	0.3	4.2
5	00ABGB	Barking and Dagenham	Eastbrook	215.23	658.46	166.35	292.75	30.18	41.68	1588.78	...	8.4	0.9	1.2	45.8	4.7	9.1

5 rows × 28 columns

5. Spatial Merge and Cluster Map

In [22]: # Load shapefile with ward boundaries
gdf = gpd.read_file("LOAC.shp")

In [23]: gdf.head()

	OA11CD	LSOA11CD	WD11CD_BF	WD11NM_BF	LAD11CD	LAD11NM	RGN11CD	LSOA11NM	USUALRES	HHOLDRES	...	HHOlds	AVHHOLDS				
0	E00023264	E01004612	E05000626	Tooting	E09000032	Wandsworth	E12000007	Wandsworth 032C		462		459	...	143		3	
1	E00003359	E01000692	E05000111	Chislehurst	E09000006	Bromley	E12000007	Bromley 002D		269		259	...	133		1	
2	E00023266	E01004615	E05000626	Tooting	E09000032	Wandsworth	E12000007	Wandsworth 034B		277		277	...	133		2	
3	E00020264	E01004027	E05000548	Riverside	E09000028	Southwark	E12000007	Southwark 003E		415		415	...	191		2	
4	E00023263	E01004613	E05000626	Tooting	E09000032	Wandsworth	E12000007	Wandsworth 033D		304		304	...	131		2	

5 rows × 22 columns

To successfully join the shapefile data with the land use dataset from the CSV, the column names containing ward names were first standardized. The shapefile originally used "WD11NM_BF" as the ward name field, while the CSV used "Area name". The column in the shapefile was renamed to ensure that both datasets shared a common key for merging.

In [24]: # Rename shapefile's ward name column to match CSV
gdf = gdf.rename(columns={"WD11NM_BF": "Area name"})

In [25]: # Perform merge on ward name
merged_gdf = gdf.merge(df_with_labels_6[['Area name', 'Cluster']], on="Area name", how="inner")

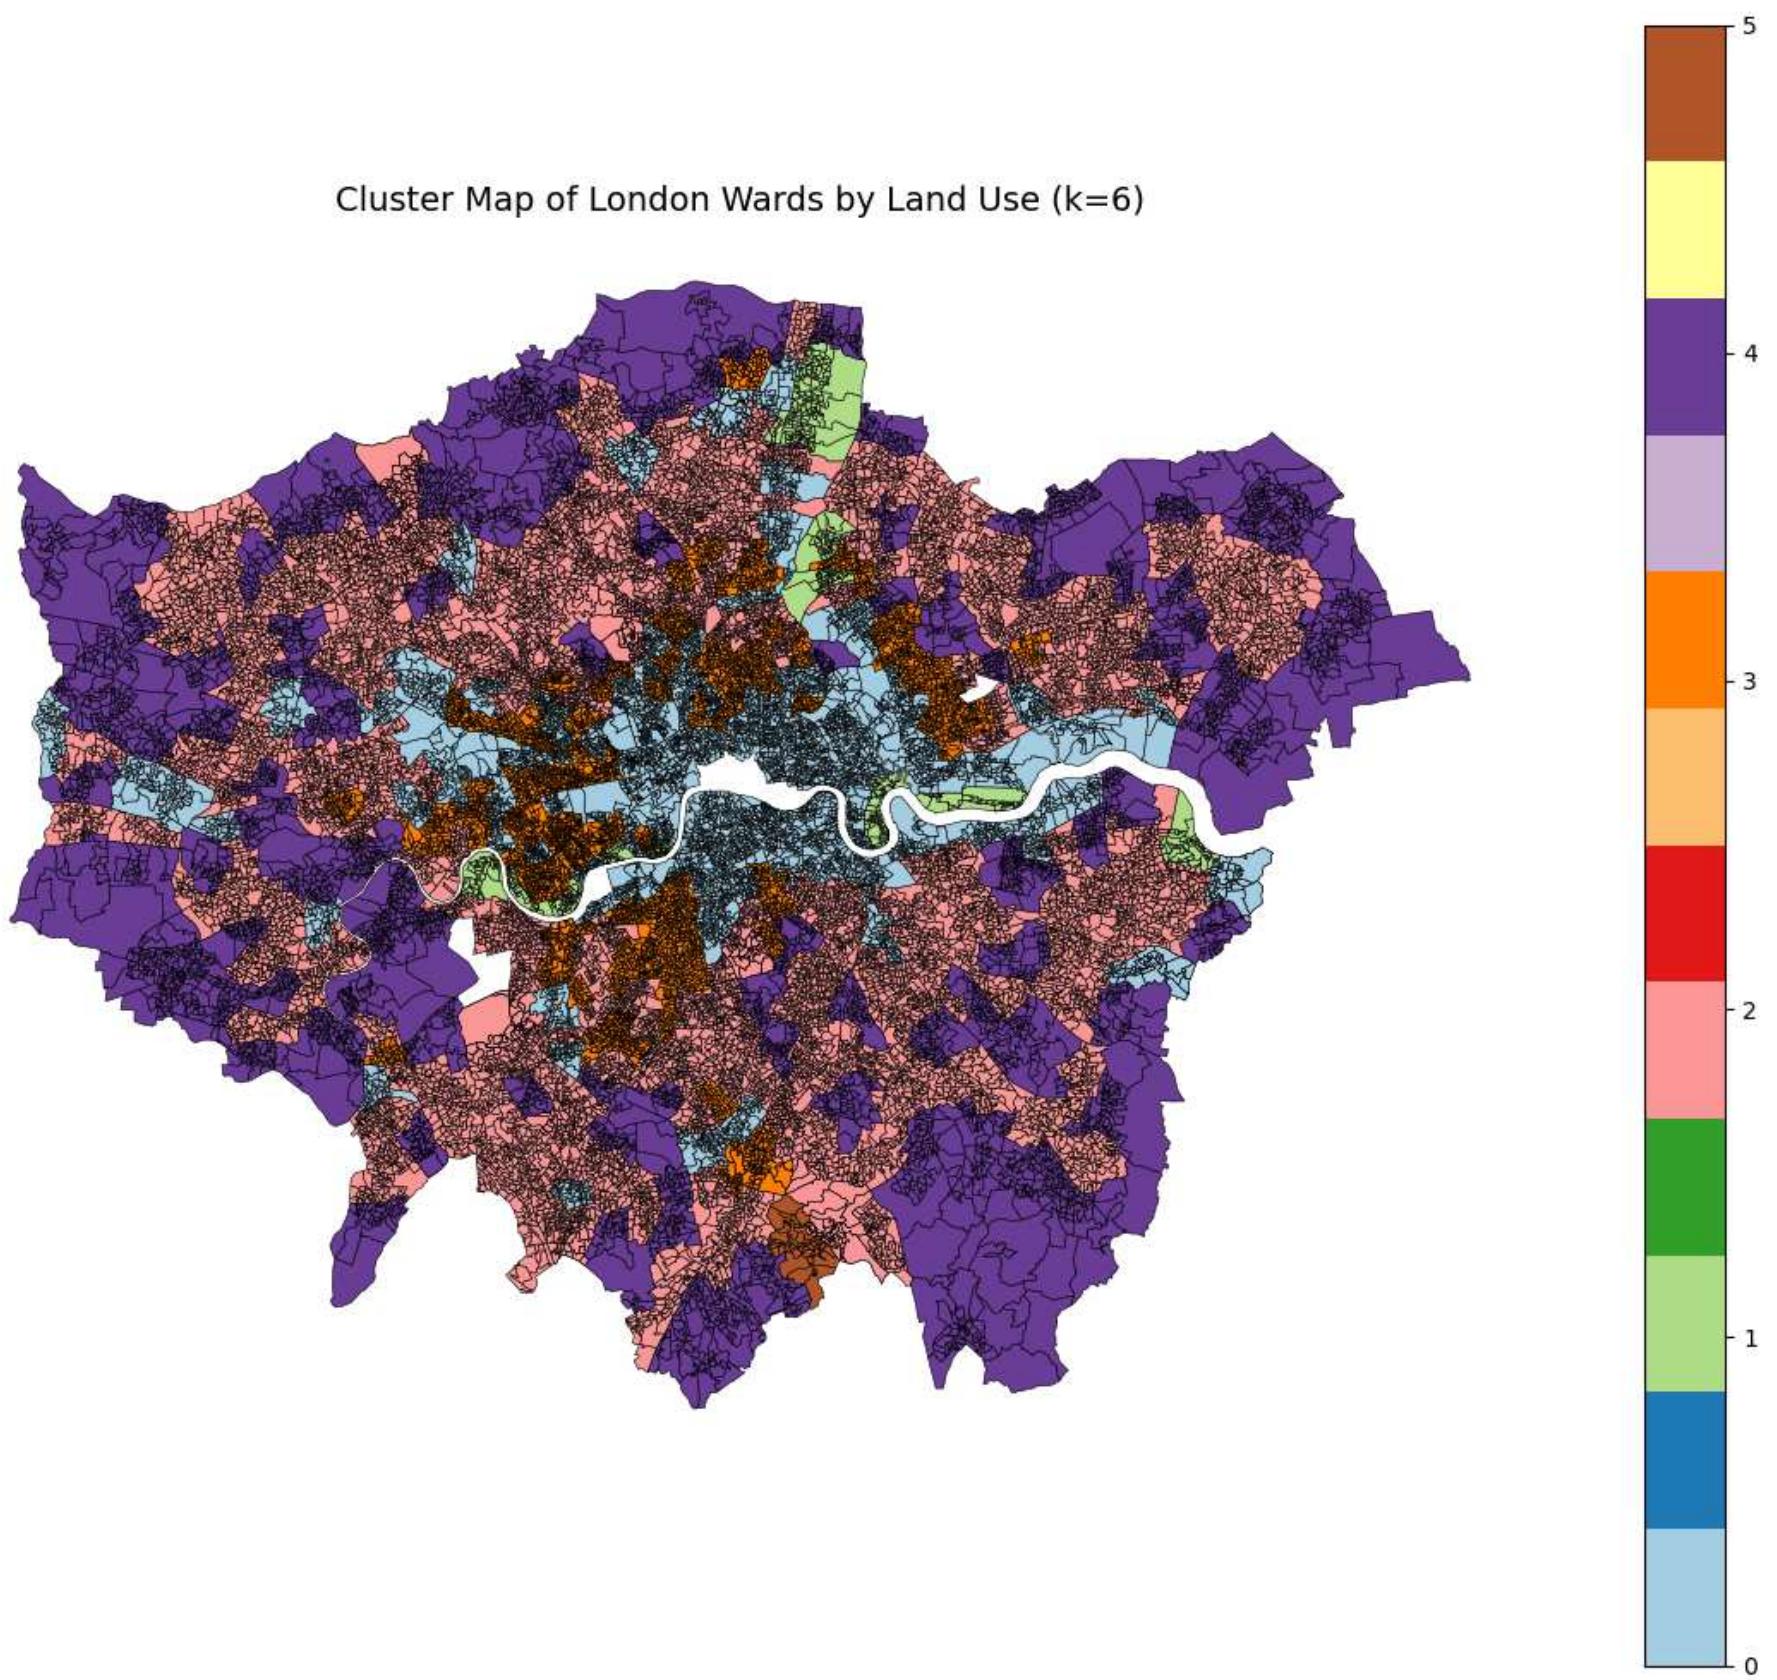
In [26]: # Cluster Map visualization
fig, ax = plt.subplots(figsize=(12, 10))
merged_gdf.plot(column="Cluster",
 cmap="Paired",
 legend=True,
 linewidth=0.4,

```

        edgecolor="black",
        ax=ax)

ax.set_title("Cluster Map of London Wards by Land Use (k=6)", fontsize=14)
ax.axis("off")
plt.tight_layout()
plt.show()

```



This map provides a spatial understanding of the 6 land use-based clusters across London's wards. Each color represents a distinct cluster characterized by a unique mix of land use types such as residential, commercial, greenspace, and transport infrastructure.

6. Cluster Profiles and Interpretation

```
In [27]: # Calculate average percentage values per cluster
cluster_features = df_with_labels_6.groupby("Cluster")[percent_columns].mean().round(2)
display(cluster_features)
```

Cluster	% Area of Domestic Buildings	% Area of Domestic Gardens	% Area of Non Domestic Buildings	% Area of Road	% Area of Rail	% Area of Path	% Area of Greenspace	% Area of Water	% Area of Other Land Uses	% Area of Unclassified Land
0	10.66	13.20	13.77	19.50	2.91	1.50	20.26	2.54	15.65	0.0
1	6.50	10.47	6.06	9.74	0.46	0.70	19.39	35.75	10.89	0.0
2	12.22	36.00	3.91	15.31	1.19	0.82	23.34	0.71	6.49	0.0
3	20.04	26.35	8.94	23.01	1.15	0.98	10.76	0.19	8.57	0.0
4	5.96	19.71	2.74	8.88	0.53	0.70	53.92	1.70	5.86	0.0
5	5.30	29.00	0.90	6.60	0.00	0.30	55.60	0.00	2.20	0.2

This table shows the average land use composition for each cluster. It provides a data-driven foundation for interpreting the spatial and functional character of each cluster, such as which ones are park-dominant, residential, transport-oriented, or mixed-use. The subsequent interpretation

section builds directly on these statistical patterns.

Summary:

Cluster Interpretation and Evaluation

Data-Driven Summary of Cluster Characteristics

The following interpretations are based on the observed average land use percentages for each cluster:

- **Cluster 0 — Mixed-Use Activity Zone**

- High % of Non-Domestic Buildings (13.77%), Roads (19.50%), and Other Land Uses (15.65%)
- Reflects transitional or commercial zones with varied land functions such as offices, industry, or large facilities.
- Moderate greenspace (20.26%) suggests limited residential or natural presence.

- **Cluster 1 — Waterfront Areas**

- Very high % Water (35.75%) with moderate Greenspace (19.39%)
- Likely includes wards adjacent to the Thames River, reservoirs, or large bodies of water.
- Built environment components are minimal, reflecting riverside or marina character.

- **Cluster 2 — Garden-Suburban Areas**

- Highest % of Domestic Gardens (36.00%) and significant Greenspace (23.34%)
- Typical of low-to-mid density suburban residential areas with private gardens and open space.
- Built environment is balanced and moderate.

- **Cluster 3 — Dense Urban Core**

- Highest % of Domestic Buildings (20.04%) and Roads (23.01%)
- Among the lowest Greenspace (10.76%)
- Clearly represents compact, inner-city zones with intense built infrastructure.

- **Cluster 4 — Greenbelt Zones**

- Very high % Greenspace (53.92%) and Domestic Gardens (19.71%)
- These wards likely lie on the urban fringe, encompassing natural areas, parks, or protected land.

- **Cluster 5 — Exceptionally Unique: Selsdon and Ballards**

- This cluster contains **only one ward: "Selsdon and Ballards"**
 - Surrounded by **Selsdon Wood**, a 200-acre National Trust woodland
 - Contains **Selsdon Park Hotel**, a historical estate formerly with a golf course
- Highest % Greenspace (55.60%) and high Domestic Gardens (29.00%)
- Possibly unique green enclaves, golf courses, or environmentally preserved areas
- Minimal built-up area suggests exceptional ecological or recreational value.

-
- K-Means clustering was applied to classify London wards into six spatially distinct urban typologies based on detailed land use compositions.
 - The clustering produced a balance between interpretability and granularity, allowing us to identify core urban, suburban, commercial, and ecological zone patterns.
 - These clusters were visualized through a spatial map and numerically summarized using average percentage metrics for key land types.

In []: